

# Homework #6

---

Student Name: Sam Crane

Student ID: 801101091

GitHub: <https://github.com/samofuture/Intro-to-ML>

```
In [ ]: # %matplotlib inline
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
import torch
import torch.optim as optim
import torch.nn as nn
# use seaborn plotting defaults
import seaborn as sns; sns.set()
print(torch.cuda.device_count())
print(f"Version: {torch.__version__}, GPU: {torch.cuda.is_available()}, NUM_GPU: {t

device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

1

Version: 2.1.1, GPU: True, NUM\_GPU: 1

## Problem 1

```
In [ ]: def prep_data() -> pd.DataFrame:
    df = pd.read_csv("Housing.csv")
    furnish_encoder = LabelEncoder()
    df['mainroad'] = df['mainroad'].apply(lambda x: 1 if x == 'yes' else 0)
    df['guestroom'] = df['guestroom'].apply(lambda x: 1 if x == 'yes' else 0)
    df['basement'] = df['basement'].apply(lambda x: 1 if x == 'yes' else 0)
    df['hotwaterheating'] = df['hotwaterheating'].apply(lambda x: 1 if x == 'yes' e
    df['airconditioning'] = df['airconditioning'].apply(lambda x: 1 if x == 'yes' e
    df['prefarea'] = df['prefarea'].apply(lambda x: 1 if x == 'yes' else 0)
    # df['furnishingstatus'] = df['furnishingstatus'].apply(lambda x: 2 if x == 'fu
    df['furnishingstatus'] = furnish_encoder.fit_transform(df['furnishingstatus'])

    return df
```

```
In [ ]: df = prep_data()
price = df.pop('price').to_numpy()
df
```

Out[ ]:

	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheati
0	7420	4	2	3	1	0	0	
1	8960	4	4	4	1	0	0	
2	9960	3	2	2	1	0	1	
3	7500	4	2	2	1	0	1	
4	7420	4	1	2	1	1	1	
...	...	...	...	...	...	...	...	
540	3000	2	1	1	1	0	1	
541	2400	3	1	1	0	0	0	
542	3620	2	1	1	1	0	0	
543	2910	3	1	1	0	0	0	
544	3850	3	1	2	1	0	0	

545 rows × 12 columns

```
In [ ]: scaler_x = StandardScaler()
data = scaler_x.fit_transform(df)

scaler_y = StandardScaler()
price = scaler_y.fit_transform(price.reshape(-1, 1))

X_train, X_test, Y_train, Y_test = train_test_split(data, price, test_size=0.2, ran

train_inputs = torch.tensor(X_train).float()
train_outputs = torch.tensor(Y_train).float()
Y_test = torch.tensor(Y_test).float()
X_test = torch.tensor(X_test).float()
```

```
In [ ]: def loss_fn(t_p, t_c):
        squared_diffs = (t_p - t_c)**2
        return squared_diffs.mean()
```

```
In [ ]: def training_loop(num_epochs, optimizer, model, t_in, t_out, v_in, v_out):
        for epoch in range(1, num_epochs+1):
            t_p = model(t_in)
            train_loss = loss_fn(t_p, t_out)
            v_p = model(v_in)
            val_loss = loss_fn(v_p, v_out)

            optimizer.zero_grad()
            train_loss.backward()
            optimizer.step()

            if epoch % 500 == 0:
```

```

        print(f"Epoch {epoch}:")
        print(f"\tTraining Loss: {float(train_loss)}")
        print(f"\tValidation Loss: {float(val_loss)}")

    return model.parameters()

```

```

In [ ]: from torch.utils.data import DataLoader, TensorDataset
        from sklearn.model_selection import KFold

        def training_loop_with_cv(num_epochs, optimizer, model, loss_fn, inputs, targets, n
            kf = KFold(n_splits=num_folds, shuffle=True, random_state=42)

            for fold, (train_idx, val_idx) in enumerate(kf.split(inputs, targets)):
                t_in, v_in = inputs[train_idx], inputs[val_idx]
                t_out, v_out = targets[train_idx], targets[val_idx]

                print(f"Training Fold {fold + 1}/{num_folds}:")

                for epoch in range(1, num_epochs + 1):
                    # Training
                    model.train()
                    t_p = model(t_in)
                    train_loss = loss_fn(t_p, t_out)

                    optimizer.zero_grad()
                    train_loss.backward()
                    optimizer.step()

                    # Validation
                    model.eval()
                    v_p = model(v_in)
                    val_loss = loss_fn(v_p, v_out)

                    if epoch % 500 == 0:
                        print(f"Epoch {epoch}:")
                        print(f"\tTraining Loss: {float(train_loss)}")
                        print(f"\tValidation Loss: {float(val_loss)}")

                return model

```

```

In [ ]: housing_model = nn.Sequential(
        nn.Linear(len(df.columns), 32),
        nn.ReLU(),
        nn.Linear(32, 1))
        optimizer = optim.SGD(housing_model.parameters(), lr=0.001)

        training_loop(5000, optimizer, housing_model, train_inputs, train_outputs, X_test,

```

```

Epoch 500:
    Training Loss: 0.43353015184402466
    Validation Loss: 0.5522839426994324
Epoch 1000:
    Training Loss: 0.34394651651382446
    Validation Loss: 0.46478649973869324
Epoch 1500:
    Training Loss: 0.31984594464302063
    Validation Loss: 0.44427594542503357
Epoch 2000:
    Training Loss: 0.30723854899406433
    Validation Loss: 0.4347897171974182
Epoch 2500:
    Training Loss: 0.29896819591522217
    Validation Loss: 0.42940303683280945
Epoch 3000:
    Training Loss: 0.2929984927177429
    Validation Loss: 0.42578548192977905
Epoch 3500:
    Training Loss: 0.28828635811805725
    Validation Loss: 0.42303478717803955
Epoch 4000:
    Training Loss: 0.28439339995384216
    Validation Loss: 0.4209675192832947
Epoch 4500:
    Training Loss: 0.2810577154159546
    Validation Loss: 0.4191916286945343
Epoch 5000:
    Training Loss: 0.2779354453086853
    Validation Loss: 0.4173828959465027

```

```
Out[ ]: <generator object Module.parameters at 0x0000023ED0791D20>
```

```

In [ ]: housing_model = nn.Sequential(
        nn.Linear(len(df.columns), 32),
        nn.ReLU(),
        nn.Linear(32, 1))
optimizer = optim.SGD(housing_model.parameters(), lr=0.001)

training_loop_with_cv(5000, optimizer, housing_model, loss_fn, train_inputs, train_

```

Training Fold 1/5:

Epoch 500:

Training Loss: 0.4466761350631714

Validation Loss: 0.44470709562301636

Epoch 1000:

Training Loss: 0.34343641996383667

Validation Loss: 0.33008602261543274

Epoch 1500:

Training Loss: 0.31954967975616455

Validation Loss: 0.3064955770969391

Epoch 2000:

Training Loss: 0.3074979782104492

Validation Loss: 0.2988772988319397

Epoch 2500:

Training Loss: 0.2990087568759918

Validation Loss: 0.29544419050216675

Epoch 3000:

Training Loss: 0.2926625907421112

Validation Loss: 0.2939257025718689

Epoch 3500:

Training Loss: 0.2873786985874176

Validation Loss: 0.2933778464794159

Epoch 4000:

Training Loss: 0.28283005952835083

Validation Loss: 0.2934132516384125

Epoch 4500:

Training Loss: 0.2788689136505127

Validation Loss: 0.2937256693840027

Epoch 5000:

Training Loss: 0.27541428804397583

Validation Loss: 0.29418814182281494

Training Fold 2/5:

Epoch 500:

Training Loss: 0.2742505371570587

Validation Loss: 0.28274160623550415

Epoch 1000:

Training Loss: 0.27066999673843384

Validation Loss: 0.28525418043136597

Epoch 1500:

Training Loss: 0.26773667335510254

Validation Loss: 0.2870347499847412

Epoch 2000:

Training Loss: 0.2650340497493744

Validation Loss: 0.2884601652622223

Epoch 2500:

Training Loss: 0.2624594271183014

Validation Loss: 0.28980302810668945

Epoch 3000:

Training Loss: 0.259947270154953

Validation Loss: 0.2909124493598938

Epoch 3500:

Training Loss: 0.2574962079524994

Validation Loss: 0.2919370234012604

Epoch 4000:

Training Loss: 0.25511255860328674

Validation Loss: 0.29308393597602844

Epoch 4500:  
    Training Loss: 0.2527482509613037  
    Validation Loss: 0.2942750155925751

Epoch 5000:  
    Training Loss: 0.2504083514213562  
    Validation Loss: 0.29550543427467346

Training Fold 3/5:

Epoch 500:  
    Training Loss: 0.233327716588974  
    Validation Loss: 0.35887157917022705

Epoch 1000:  
    Training Loss: 0.22994662821292877  
    Validation Loss: 0.36882466077804565

Epoch 1500:  
    Training Loss: 0.22731490433216095  
    Validation Loss: 0.3757067322731018

Epoch 2000:  
    Training Loss: 0.22506628930568695  
    Validation Loss: 0.3808686137199402

Epoch 2500:  
    Training Loss: 0.22309818863868713  
    Validation Loss: 0.38479289412498474

Epoch 3000:  
    Training Loss: 0.2212824523448944  
    Validation Loss: 0.38808673620224

Epoch 3500:  
    Training Loss: 0.21962371468544006  
    Validation Loss: 0.39074578881263733

Epoch 4000:  
    Training Loss: 0.2180585414171219  
    Validation Loss: 0.39309006929397583

Epoch 4500:  
    Training Loss: 0.21660666167736053  
    Validation Loss: 0.3947078585624695

Epoch 5000:  
    Training Loss: 0.21524566411972046  
    Validation Loss: 0.3964630961418152

Training Fold 4/5:

Epoch 500:  
    Training Loss: 0.25040555000305176  
    Validation Loss: 0.2273048311471939

Epoch 1000:  
    Training Loss: 0.24664267897605896  
    Validation Loss: 0.2334004044532776

Epoch 1500:  
    Training Loss: 0.24386152625083923  
    Validation Loss: 0.23786869645118713

Epoch 2000:  
    Training Loss: 0.24140538275241852  
    Validation Loss: 0.24157866835594177

Epoch 2500:  
    Training Loss: 0.2392292469739914  
    Validation Loss: 0.24478623270988464

Epoch 3000:  
    Training Loss: 0.23725339770317078  
    Validation Loss: 0.2477830946445465

```

Epoch 3500:
    Training Loss: 0.23546332120895386
    Validation Loss: 0.25030621886253357
Epoch 4000:
    Training Loss: 0.23370973765850067
    Validation Loss: 0.25262802839279175
Epoch 4500:
    Training Loss: 0.23196786642074585
    Validation Loss: 0.2549841105937958
Epoch 5000:
    Training Loss: 0.23034237325191498
    Validation Loss: 0.2569766938686371
Training Fold 5/5:
Epoch 500:
    Training Loss: 0.2502225339412689
    Validation Loss: 0.1641637533903122
Epoch 1000:
    Training Loss: 0.24779418110847473
    Validation Loss: 0.16780084371566772
Epoch 1500:
    Training Loss: 0.24593347311019897
    Validation Loss: 0.17067712545394897
Epoch 2000:
    Training Loss: 0.24424734711647034
    Validation Loss: 0.17314819991588593
Epoch 2500:
    Training Loss: 0.24266591668128967
    Validation Loss: 0.17531979084014893
Epoch 3000:
    Training Loss: 0.24117106199264526
    Validation Loss: 0.1774040013551712
Epoch 3500:
    Training Loss: 0.23971763253211975
    Validation Loss: 0.17939984798431396
Epoch 4000:
    Training Loss: 0.23827409744262695
    Validation Loss: 0.18126295506954193
Epoch 4500:
    Training Loss: 0.23689046502113342
    Validation Loss: 0.18294687569141388
Epoch 5000:
    Training Loss: 0.2355634868144989
    Validation Loss: 0.18456332385540009

```

```

Out[ ]: Sequential(
  (0): Linear(in_features=12, out_features=32, bias=True)
  (1): ReLU()
  (2): Linear(in_features=32, out_features=1, bias=True)
)

```

The model complexity is orders of magnitude higher than linear regression. With each neuron in a neural network essentially being a linear regression on its own, this is a much more complex task. When we increased the number of layers for part B, the complexity only increased.

Compared to HW5, the first network performed worse than the linear regression. However, the second model was performing just as well, if not better than HW5.

## Problem 2

```
In [ ]: from torchvision import datasets, transforms
data_path = 'data-unversioned/p1ch7'
cifar10 = datasets.CIFAR10(data_path, train=True, download=True)
cifar10_val = datasets.CIFAR10(data_path, train=False, download=True)
```


c:\Users\Sam\anaconda3\Lib\site-packages\torchvision\io\image.py:13: UserWarning: Failed to load image Python extension: 'Could not find module 'C:\Users\Sam\anaconda3\Lib\site-packages\torchvision\image.pyd' (or one of its dependencies). Try using the full path with constructor syntax.'If you don't plan on using image functionality from `torchvision.io`, you can ignore this warning. Otherwise, there might be something wrong with your environment. Did you have `libjpeg` or `libpng` installed before building `torchvision` from source?

warn(

Files already downloaded and verified

Files already downloaded and verified

```
In [ ]: img, label = cifar10[99]
plt.imshow(img)
plt.show()
```



```
In [ ]: to_tensor = transforms.ToTensor()

cifar10_tensor = datasets.CIFAR10(data_path, train=True, download=False, transform=
```



```

imgs = torch.stack([img_t for img_t, _ in cifar10_tensor], dim=3)
mean = imgs.view(3, -1).mean(dim=1)
std = imgs.view(3, -1).std(dim=1)

# print("M", type(mean))
# print("S", std)
# transforms.Normalize(mean, std)

```

```

In [ ]: cifar10 = datasets.CIFAR10(
        data_path, train=True, download=False,
        transform=transforms.Compose([
            transforms.ToTensor(),
            transforms.Normalize(mean, std)
        ]))

# cifar10 = [(img, label) for img, label in cifar10]

cifar10_val = datasets.CIFAR10(
    data_path, train=False, download=False,
    transform=transforms.Compose([
        transforms.ToTensor(),
        transforms.Normalize(mean, std)
    ]))

# cifar10_val = [(img, label) for img, label in cifar10_val]

```

```

In [ ]: class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
                       'dog', 'frog', 'horse', 'ship', 'truck']

n_out = 10

cifar_model = nn.Sequential(
    nn.Linear(3072, 512),
    nn.ReLU(),
    nn.Linear(512, n_out),
    nn.Softmax(dim=1)
)

loss_fn = nn.CrossEntropyLoss()
optimizer = optim.SGD(cifar_model.parameters(), lr = 0.01)
num_epochs = 50

```

```

In [ ]: cifar_model.to(device)
        loss_fn.to(device)

        for epoch in range(num_epochs):
            for img, label in cifar10:
                img = torch.Tensor(img)
                img = img.view(-1).unsqueeze(0)
                label = torch.tensor([label])
                img = img.to(device)
                out = cifar_model(img)
                loss = loss_fn(out, label.to(device))

```

```
optimizer.zero_grad()  
loss.backward()  
optimizer.step()  
  
print("Epoch: %d, Loss: %f" % (epoch, float(loss)))
```

```
Epoch: 0, Loss: 2.461141  
Epoch: 1, Loss: 1.536038  
Epoch: 2, Loss: 2.454097  
Epoch: 3, Loss: 2.461150  
Epoch: 4, Loss: 2.461084  
Epoch: 5, Loss: 2.461150  
Epoch: 6, Loss: 2.461150  
Epoch: 7, Loss: 2.461150  
Epoch: 8, Loss: 2.461150  
Epoch: 9, Loss: 2.461150  
Epoch: 10, Loss: 2.461150  
Epoch: 11, Loss: 2.461150  
Epoch: 12, Loss: 2.461150  
Epoch: 13, Loss: 2.456374  
Epoch: 14, Loss: 2.461150  
Epoch: 15, Loss: 2.461150  
Epoch: 16, Loss: 2.461150  
Epoch: 17, Loss: 2.460614  
Epoch: 18, Loss: 2.461150  
Epoch: 19, Loss: 2.461150  
Epoch: 20, Loss: 2.461150  
Epoch: 21, Loss: 2.461150  
Epoch: 22, Loss: 2.461150  
Epoch: 23, Loss: 2.461150  
Epoch: 24, Loss: 2.461150  
Epoch: 25, Loss: 2.460831  
Epoch: 26, Loss: 2.461150  
Epoch: 27, Loss: 2.461150  
Epoch: 28, Loss: 2.461150  
Epoch: 29, Loss: 2.461150  
Epoch: 30, Loss: 2.461150  
Epoch: 31, Loss: 2.461150  
Epoch: 32, Loss: 2.461150  
Epoch: 33, Loss: 2.461150  
Epoch: 34, Loss: 2.461150  
Epoch: 35, Loss: 2.461150  
Epoch: 36, Loss: 2.461150  
Epoch: 37, Loss: 2.461150  
Epoch: 38, Loss: 2.461150  
Epoch: 39, Loss: 2.461150  
Epoch: 40, Loss: 2.461150  
Epoch: 41, Loss: 2.461150  
Epoch: 42, Loss: 2.461150  
Epoch: 43, Loss: 2.461150  
Epoch: 44, Loss: 2.461150  
Epoch: 45, Loss: 2.461150  
Epoch: 46, Loss: 2.461150  
Epoch: 47, Loss: 2.461150  
Epoch: 48, Loss: 2.461150  
Epoch: 49, Loss: 2.461150
```

```
In [ ]: train_loader = torch.utils.data.DataLoader(cifar10_val, batch_size=64, shuffle=False)

correct = 0
total = 0
cifar_model.to(device)
loss_fn.to(device)
with torch.no_grad():
    for imgs, labels in train_loader:
        imgs = torch.Tensor(imgs).to(device)

        labels = torch.tensor(labels).to(device)

        outputs = cifar_model(imgs.view(imgs.shape[0], -1))
        _, predicted = torch.max(outputs, dim=1)
        total += labels.shape[0]
        correct += int((predicted == labels).sum())

print("Accuracy: %f" % (correct / total))
```

C:\Users\Sam\AppData\Local\Temp\ipykernel\_35472\3286169161.py:11: UserWarning: To copy construct from a tensor, it is recommended to use sourceTensor.clone().detach() or sourceTensor.clone().detach().requires\_grad\_(True), rather than torch.tensor(sourceTensor).

```
labels = torch.tensor(labels).to(device)
```

Accuracy: 0.365300

```
In [ ]: import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 16, kernel_size=3, padding=1)
        self.conv2 = nn.Conv2d(16, 8, kernel_size=3, padding=1)
        self.fc1 = nn.Linear(8 * 8 * 8, 32)
        self.fc2 = nn.Linear(32, 10)

    def forward(self, x):
        out = F.max_pool2d(torch.tanh(self.conv1(x)), 2)
        out = F.max_pool2d(torch.tanh(self.conv2(out)), 2)
        out = out.view(-1, 8 * 8 * 8)
        out = torch.tanh(self.fc1(out))
        out = self.fc2(out)
        return out
```

```
In [ ]: model = nn.Sequential(
    nn.Linear(3072, 512),
    nn.ReLU(),
    nn.Linear(512, 256),
    nn.ReLU(),
    nn.Linear(256, 128),
    nn.ReLU(),
    nn.Linear(128, n_out),
    nn.Softmax(dim=1)
)
```

```
In [ ]: import datetime # <1>

def training_loop(n_epochs, optimizer, model, loss_fn, train_loader):
    for epoch in range(1, n_epochs + 1): # <2>
        loss_train = 0.0
        for imgs, labels in train_loader: # <3>

            outputs = model(imgs) # <4>

            loss = loss_fn(outputs, labels) # <5>

            optimizer.zero_grad() # <6>

            loss.backward() # <7>

            optimizer.step() # <8>

            loss_train += loss.item() # <9>

        if epoch == 1 or epoch % 10 == 0:
            print('{} Epoch {}, Training loss {}'.format(
                datetime.datetime.now(), epoch,
                loss_train / len(train_loader)))
```

```
In [ ]: train_loader = torch.utils.data.DataLoader(cifar10, batch_size=64, shuffle=True)

model = Net() # <2>
optimizer = optim.SGD(model.parameters(), lr=1e-2) # <3>
loss_fn = nn.CrossEntropyLoss() # <4>

training_loop( # <5>
    n_epochs = 100,
    optimizer = optimizer,
    model = model,
    loss_fn = loss_fn,
    train_loader = train_loader,
)
```

```
2023-12-02 18:39:31.725474 Epoch 1, Training loss 2.0475708695933643
2023-12-02 18:42:43.357330 Epoch 10, Training loss 1.1856087433255238
2023-12-02 18:46:17.277680 Epoch 20, Training loss 1.0059958333554475
2023-12-02 18:49:52.609951 Epoch 30, Training loss 0.9336449801921844
2023-12-02 18:53:28.861767 Epoch 40, Training loss 0.8844456407420166
2023-12-02 18:57:01.941340 Epoch 50, Training loss 0.8441681743354139
2023-12-02 19:00:42.824281 Epoch 60, Training loss 0.8128206913199876
2023-12-02 19:04:21.251485 Epoch 70, Training loss 0.7862183522156743
2023-12-02 19:07:59.140729 Epoch 80, Training loss 0.7612364073391156
2023-12-02 19:11:38.193575 Epoch 90, Training loss 0.7391753197478517
2023-12-02 19:15:18.577279 Epoch 100, Training loss 0.7216365825565879
```

```
In [ ]: def validate(model, loader):
    correct = 0
    total = 0

    with torch.no_grad(): # <1>
```

```
for imgs, labels in loader:
    outputs = model(imgs)
    _, predicted = torch.max(outputs, dim=1) # <2>
    total += labels.shape[0] # <3>
    correct += int((predicted == labels).sum()) # <4>

print("Accuracy: {:.2f}".format(correct / total))

val_loader = torch.utils.data.DataLoader(cifar10_val, batch_size=64, shuffle=False)
validate(model, val_loader)
```

Accuracy: 0.65

## Part A

Training Time: 2 Hours Training Loss: 2.5144 Evaluation Accuracy: 0.36

## Part B

Training Time: 3 Hours Training Loss: 0.72 Evaluation Accuracy: 0.65

I don't think the model trained long enough to achieve overfitting. The complexity of the network adds  $128 + 256 = 384$  more neurons to the already existing 522, making it more than 50% more complex.